

Data Science Challenge 2021 - Jane Holland

This kernel aims to predict car insurance cold call success. It shows data exploration, feature engineering, visualization and different classifier selection.

The kernel is split into 6 main sections:

1. Viewing the data - what are we working with.
2. Cleaning the data - preprocessing the data we are working with.
3. Data Analysis and Visualisation - what trends and patterns we can see in the data.
4. Prepare Data for Machine Learning - change category columns into numerical data.
5. Machine Learning Classifiers - using different machine learning techniques to predict the cold call success
6. Testing the Trained Model - examine the model on the test data (CarInsurance_test.csv)

Import necessary libraries

```
In [48]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import operator
import itertools
from itertools import product

from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder

# Regression & Classification Models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier, RandomForestClassifier
from xgboost import XGBClassifier
from sklearn import tree

# Metrics
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

sns.set_theme(style="whitegrid")
```

1. Viewing the data

Read in csv files and show first 5 rows

```
In [49]: df_train = pd.read_csv('carInsurance_train.csv', sep=',')
df_test = pd.read_csv('carInsurance_test.csv', sep=',')
```

```
df_train.head()
```

```
Out[49]:
```

		Id	Age	Job	Marital	Education	Default	Balance	HHInsurance	CarLoan	Communicatio
0	1	32	management	single	tertiary	0	1218	1	0	telephon	
1	2	32	blue-collar	married	primary	0	1156	1	0	Nal	
2	3	29	management	single	tertiary	0	637	1	0	cellula	
3	4	25	student	single	primary	0	373	1	0	cellula	
4	5	30	management	married	tertiary	0	2694	0	0	cellula	

Display the number of rows and columns as a tuple

```
In [50]: df_train.shape
```

```
Out[50]: (4000, 19)
```

Display information about dataframe; columns, datatypes, etc.

```
In [51]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    4000 non-null   int64
1   Age                  4000 non-null   int64
2   Job                  3981 non-null   object
3   Marital              4000 non-null   object
4   Education            3831 non-null   object
5   Default              4000 non-null   int64
6   Balance              4000 non-null   int64
7   HHInsurance          4000 non-null   int64
8   CarLoan              4000 non-null   int64
9   Communication        3098 non-null   object
10  LastContactDay       4000 non-null   int64
11  LastContactMonth     4000 non-null   object
12  NoOfContacts         4000 non-null   int64
13  DaysPassed           4000 non-null   int64
14  PrevAttempts         4000 non-null   int64
15  Outcome              958 non-null    object
16  CallStart            4000 non-null   object
17  CallEnd              4000 non-null   object
18  CarInsurance         4000 non-null   int64
dtypes: int64(11), object(8)
memory usage: 593.9+ KB
```

2. Cleaning the data (preprocessing)

This involves:

- removing any data that is not considered important.

- filling in missing data.
- checking for duplicates (and removing if present).
- converting datatypes to the correct datatype.
- dealing with outliers.

Show any null values in dataframe

We will need to clean any NA values in the dataframe by either 1) dropping the column, or 2) filling in the missing data with the mode or the mean.

```
In [52]: df_train.isnull().sum()
```

```
Out[52]: Id                0
Age                0
Job                19
Marital            0
Education          169
Default            0
Balance            0
HHInsurance        0
CarLoan            0
Communication      902
LastContactDay     0
LastContactMonth   0
NoOfContacts       0
DaysPassed         0
PrevAttempts       0
Outcome            3042
CallStart          0
CallEnd            0
CarInsurance       0
dtype: int64
```

Show any duplicates in dataframe

```
In [53]: duplicates = df_train.duplicated()
print(duplicates.sum())
```

```
0
```

Remove and fill columns

- As we are not conducting any time series on individuals, the Id column can be removed.
- The Outcome column has 3042/4000 missing values, so this column can also be dropped.
- The Job, Education, and Communication columns are missing only a small amount of data so we will fill them with the most common attributes.

```
In [54]: # Drop ID (not sequential, independent) and Outcome (too much missing data) columns
df_train = df_train.drop(['Id'], axis=1)
df_train = df_train.drop(['Outcome'], axis=1)

# Fill NA values with the most common attribute for each column
for column in ['Job', 'Education', 'Communication']:
    df_train[column] = df_train[column].fillna(df_train[column].mode()[0])
```

```
df_train.isnull().sum()
```

```
Out[54]: Age          0
Job          0
Marital      0
Education    0
Default      0
Balance      0
HHInsurance  0
CarLoan      0
Communication 0
LastContactDay 0
LastContactMonth 0
NoOfContacts 0
DaysPassed   0
PrevAttempts 0
CallStart    0
CallEnd      0
CarInsurance 0
dtype: int64
```

Converting datatypes

The CallStart and CallEnd need to be changed from objects to the time format.

```
In [55]: # Conver to time datatype
df_train[['CallStart', 'CallEnd']] = df_train[['CallStart', 'CallEnd']].astype('datetime64[ns]')

# While we are dealing with the time attributes we will create a new column call duration

#Total Call Duration
df_train['CallDuration'] = df_train['CallEnd'] - df_train['CallStart']

# To make the CallDuration more comprehensible, we will convert it to seconds (or minutes)
df_train['CallDuration'] = df_train['CallDuration'].dt.components['minutes'] * 60 + df_train['CallDuration'].dt.components['seconds']
```

Taking a closer look at the string objects in dataframe

```
In [56]: df_train.describe(include=['O'])
```

```
Out[56]:
```

	Job	Marital	Education	Communication	LastContactMonth
count	4000	4000	4000	4000	4000
unique	11	3	3	2	12
top	management	married	secondary	cellular	may
freq	912	2304	2157	3733	1049

Do the same with the numerical datatypes

```
In [57]: df_train.describe(datetime_is_numeric=True)
```

```
Out[57]:
```


	Age	Default	Balance	HHInsurance	CarLoan	LastContactDay	NoOfContacts
count	4000.000000	4000.000000	4000.000000	4000.000000	4000.000000	4000.000000	4000.000000

	Age	Default	Balance	HHInsurance	CarLoan	LastContactDay	NoOfCont
mean	41.214750	0.014500	1532.937250	0.49275	0.133000	15.721250	2.60
min	18.000000	0.000000	-3058.000000	0.00000	0.000000	1.000000	1.00
25%	32.000000	0.000000	111.000000	0.00000	0.000000	8.000000	1.00
50%	39.000000	0.000000	551.500000	0.00000	0.000000	16.000000	2.00
75%	49.000000	0.000000	1619.000000	1.00000	0.000000	22.000000	3.00
max	95.000000	1.000000	98417.000000	1.00000	1.000000	31.000000	43.00
std	11.550194	0.119555	3511.452489	0.50001	0.339617	8.425307	3.06

Handling the outliers

In the above description we can see that the mean of the balance is 1532, but the min value is -3058 and the max value is 98417. This seems like a large range so we will take a closer look and use a boxplot to see what the spread is like.

```
In [58]: plt.figure(figsize=(16, 6))
ax = sns.boxplot(data=df_train, orient="h", palette="Set2")
```



As suspected the Balance column has some outliers. we will create a function to remove the outliers and only use the values between the 1st and 3rd quartile

```
In [59]: # remove outliers in Balance column
def remove_outliers(col):
    sorted(col)
    Q1,Q3=col.quantile([0.25,0.75])
    # Use interquartile range (IQR) to find outliers.
    IQR=Q3-Q1
    # Values lower than the lower boundary are outliers.
    lower_range=Q1-(1.5*IQR)
    # Values greater than the upper boundary are outliers.
    upper_range=Q3+(1.5*IQR)
    return lower_range, upper_range
```

```
In [60]: # Redefining the Balance column - removing the outliers
low_bal, upp_bal = remove_outliers(df_train['Balance'])
df_train['Balance']=np.where(df_train['Balance'] > upp_bal, upp_bal, df_train['Balance'])
df_train['Balance']=np.where(df_train['Balance'] < low_bal, low_bal, df_train['Balance'])
```

3. Data Analysis and Visualisation

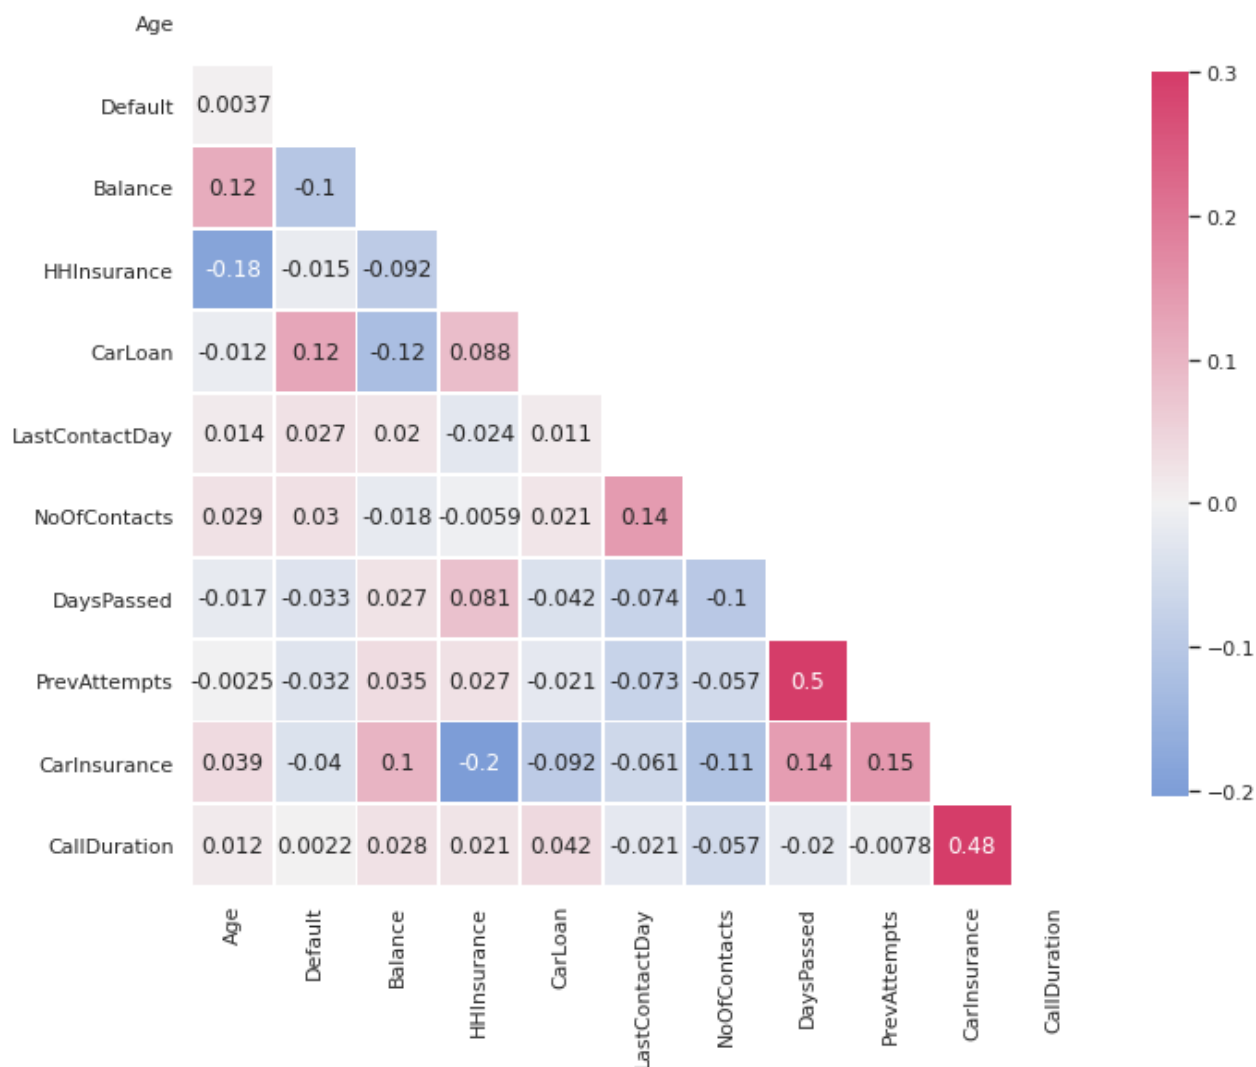
Here we will explore the data in more detail and visualise any significant observations.

Data Correlation

```
In [61]: # Display the correlated data as a heatmap - this will show the positive and negative correlations
corr = df_train.corr()
mask = np.triu(np.ones_like(corr, dtype=bool))
f, ax = plt.subplots(figsize=(11, 9))
cmap = sns.diverging_palette(255, 0, as_cmap=True)

sns.heatmap(corr, annot=True, fmt=".2", mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.8, cbar_kws={"shrink": .8})
```

Out[61]: <matplotlib.axes._subplots.AxesSubplot at 0x7f19fcd9190>



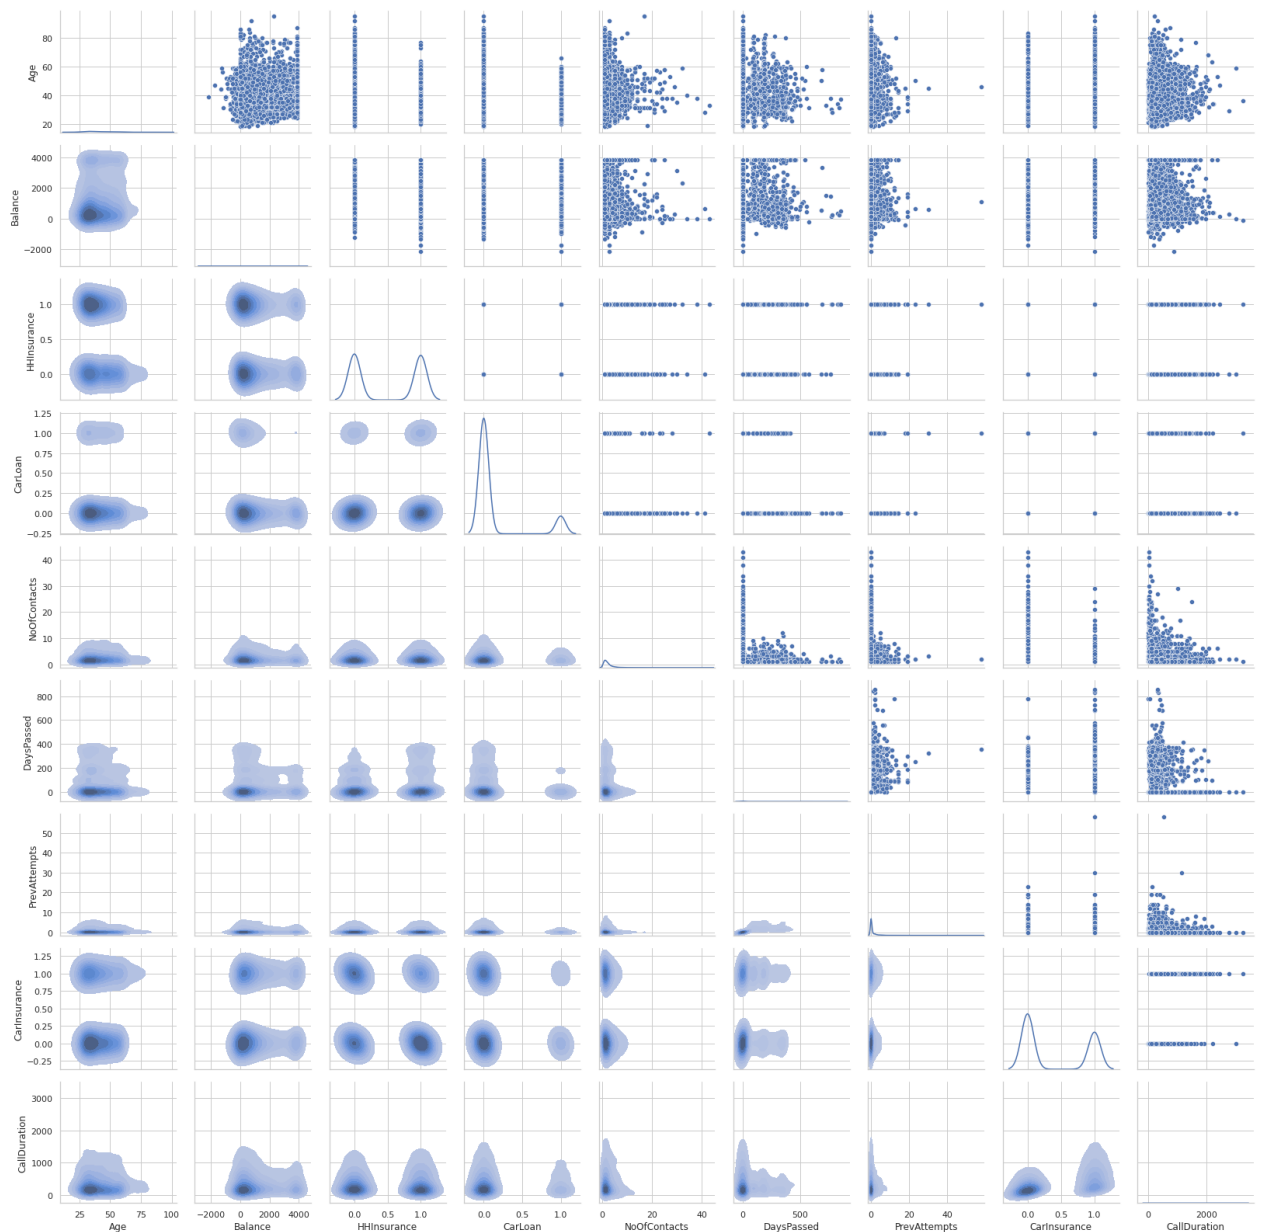
From this heatmap correlation matrix we can see that:

- There are Strong Positive correlations between the CarInsurance and CallDuration, as well as PrevAttempts and DaysPassed. Negative correlations can be seen with the LastContactDay and CarInsurance, as well as the Age and HHInsurance.
- Perhaps a longer CallDuration increases the odds of subscribing car insurance.
- Furthermore, the more times an individual is contacted, the less likely they are to subscribe to car insurance.
- Other attributes are fairly independent.

Looking more closely at the pairwise relationships in a dataset

```
In [62]: # Plot some of the more important attribute features  
df_pair_feats = ['Age', 'Balance', 'HHInsurance', 'CarLoan', 'NoOfContacts', 'DaysPa  
g = sns.PairGrid(df_train[df_pair_feats])  
g.map_upper(sns.scatterplot)  
g.map_lower(sns.kdeplot, fill=True)  
g.map_diag(sns.kdeplot)  
g.add_legend()
```

```
Out[62]: <seaborn.axisgrid.PairGrid at 0x7f19f53ed280>
```



Here we can see a few different observations:

- Younger people (<40) are less likely to have a car loan or car insurance subscribed. They are more likely to have house insurance and have a smaller balance.
- People with house insurance are less likely to have car insurance subscribed.
- The longer the Call Duration the more likely car insurance will be subscribed.
- The more frequently the number of times contacted during this campaign, the less likely that car insurance will be subscribed to the individual.
- The more previous attempts in the previous campaign, the more likely of subscribing insurance.

Category Analysis and Visualisation

```
In [63]: # Display the number and percentage of jobs that have car insurance subscribed
job_insured=pd.crosstab(df_train['Job'],df_train['CarInsurance'],colnames=['Car
```



```
job_insured['% with Car Insurance']=round(job_insured[1]/(job_insured[0]+job_insured[1]))
job_insured
```

Out[63]:

Car Insurance	0	1	% with Car Insurance
Job			
admin.	274	185	40.31
blue-collar	540	219	28.85
entrepreneur	86	35	28.93
housemaid	72	37	33.94
management	511	401	43.97
retired	103	146	58.63
self-employed	86	54	38.57
services	218	112	33.94
student	44	87	66.41
technician	406	254	38.48
unemployed	56	74	56.92

In [64]:

```
# Display the number and percentage of different marital statuses that have car insurance
marital_insured=pd.crosstab(df_train['Marital'],df_train['CarInsurance'],colname='% with Car Insurance')
marital_insured['% with Car Insurance']=round(marital_insured[1]/(marital_insured[0]+marital_insured[1]))
marital_insured
```




Out[64]:


Car Insurance	0	1	% with Car Insurance
Marital			
divorced	273	210	43.48
married	1471	833	36.15
single	652	561	46.25

In [65]:

```
# Plot the amount of category features that have car insurance subscribed
df_category_feats = df_train[['Job','Marital','Education','Communication']]


for feature in df_category_feats:
    plt.figure(figsize=(10,4))
    sns.barplot(x=feature, y='CarInsurance', data= df_train, palette='magma')
    plt.xticks(rotation=45)
    plt.ylabel('Car Insurance')
    plt.tight_layout()
```






```
In [66]: # Plotting the previous campaign (PrevAttempts) vs the current campaign (NoOfCor
# Plot how the number of previous attempts affects the outcome of car insurance
plt.figure(figsize=(7,5))
sns.barplot( x='CarInsurance', y='PrevAttempts', data=df_train, palette='magma')
plt.ylabel('No of Previous Attempts')
plt.xlabel('Car Insurance')
plt.show()

# Plot how the number of contacts affects the outcome of car insurance subscrib
plt.figure(figsize=(7,5))
sns.barplot( x='CarInsurance', y='NoOfContacts', data=df_train, palette='magma')
plt.ylabel('No of Contacts')
plt.xlabel('Car Insurance')
plt.show()
```





Feature Engineering

By using feature engineering we can create more features that can be further analysed as well as for grouping attributes.

The Balance, Age, LastContactDay, and CallStart attributes are good examples of attributes to discretise it into equal-sized buckets. This produces a Category object where each quantile or bracket has a membership for each datapoint.

```
In [67]: # Here I will create the bins and give each one a label.
# For e.g. if your balance -500 you will be in the 'Negative Balance' bracket

# Whats the min and max balances - this can inform decision on how many bins to
print("Train:", df_train['Balance'].min(), df_train['Balance'].max())
print("Test:", df_test['Balance'].min(), df_test['Balance'].max()) # outliers he

# Create and label bins
train_bal_bins= [-3000,0,1295,2590,3885]
train_bal_labels = ['Negative Balance','Low','Mid','High']

# Create a new column 'BalanceBracket' that shows the balance groups
df_train['BalanceBracket'] = pd.cut(df_train['Balance'], bins=train_bal_bins, la

# Display the number and percentage of different balance brackets that have car
balance_insured=pd.crosstab(df_train['BalanceBracket'],df_train['CarInsurance'],
balance_insured['% with Car Insurance']=round(balance_insured[1]/(balance_insured
balance_insured
```

Train: -2151.0 3881.0

Test: -1980 41630


```
Out[67]: Car Insurance    0    1  % with Car Insurance
```

BalanceBracket			
Negative Balance	207	58	21.89
Low	1571	995	38.78

Car Insurance	0	1	% with Car Insurance
BalanceBracket			
Mid	270	245	47.57
High	348	306	46.79

In [68]: *# Plotting how the call duration affects the likelihood of subscribing to car insurance for different balance groups*

```
plt.figure(figsize=(10,6))
sns.barplot(x='BalanceBracket', y='CallDuration', hue='CarInsurance', data=df_train)
plt.ylabel('Call Duration (secs)')
plt.xlabel('Balance Demographics')
plt.show()
```



In [69]: *# For the age brackets I chose intervals of 10 in order to make the data more continuous*

```
# Create the bins and give each one a label.
train_age_bins = [18,21,31,41,51,61,71,81,91,101]
train_age_labels = ['18-21','21-30','31-40','41-50','51-60','61-70','71-80','81-90','91-101']

# Create a new column 'AgeBracket' that shows the age groups
df_train['AgeBracket'] = pd.cut(df_train['Age'], bins=train_age_bins, labels=train_age_labels)

# Display the number and percentage of different age brackets that have car insurance
age_insured = pd.crosstab(df_train['AgeBracket'], df_train['CarInsurance'], colnames=['AgeBracket', 'Car Insurance'])
age_insured['% with Car Insurance'] = round(age_insured[1]/(age_insured[0]+age_insured[1]), 2)
age_insured
```


Out[69]:

Car Insurance	0	1	% with Car Insurance
AgeBracket			
18-21	7	16	69.57

Car Insurance	0	1	% with Car Insurance
AgeBracket			
21-30	342	313	47.79
31-40	955	544	36.29
41-50	618	331	34.88
51-60	434	250	36.55
61-70	21	95	81.90
71-80	16	45	73.77
81-90	3	8	72.73
90+	0	2	100.00

In [70]: *# Plotting how the call duration affects the likelihood of subscribing to car insurance for different age demographics*

```
plt.figure(figsize=(10,6))
sns.barplot(x='AgeBracket', y='CallDuration', hue='CarInsurance', data=df_train,
plt.ylabel('Call Duration (secs)')
plt.xlabel('Age Demographics')
plt.show()
```



In [71]: *# For the LastContactDay brackets I chose every 11 days to make the data more clear*

```
# Create the bins and give each one a label.
train_day_bins= [0,10,21,32]
train_day_labels = ['Start','Middle','End']

# Create a new column 'DayBracket' that shows the age groups
df_train['DayBracket'] = pd.cut(df_train['LastContactDay'], bins=train_day_bins,
```

```
# Display the number and percentage of different contact day brackets that have
day_insured=pd.crosstab(df_train['DayBracket'],df_train['CarInsurance'],colnames
day_insured['% with Car Insurance']=round(day_insured[1]/(day_insured[0]+day_ins
day_insured
```

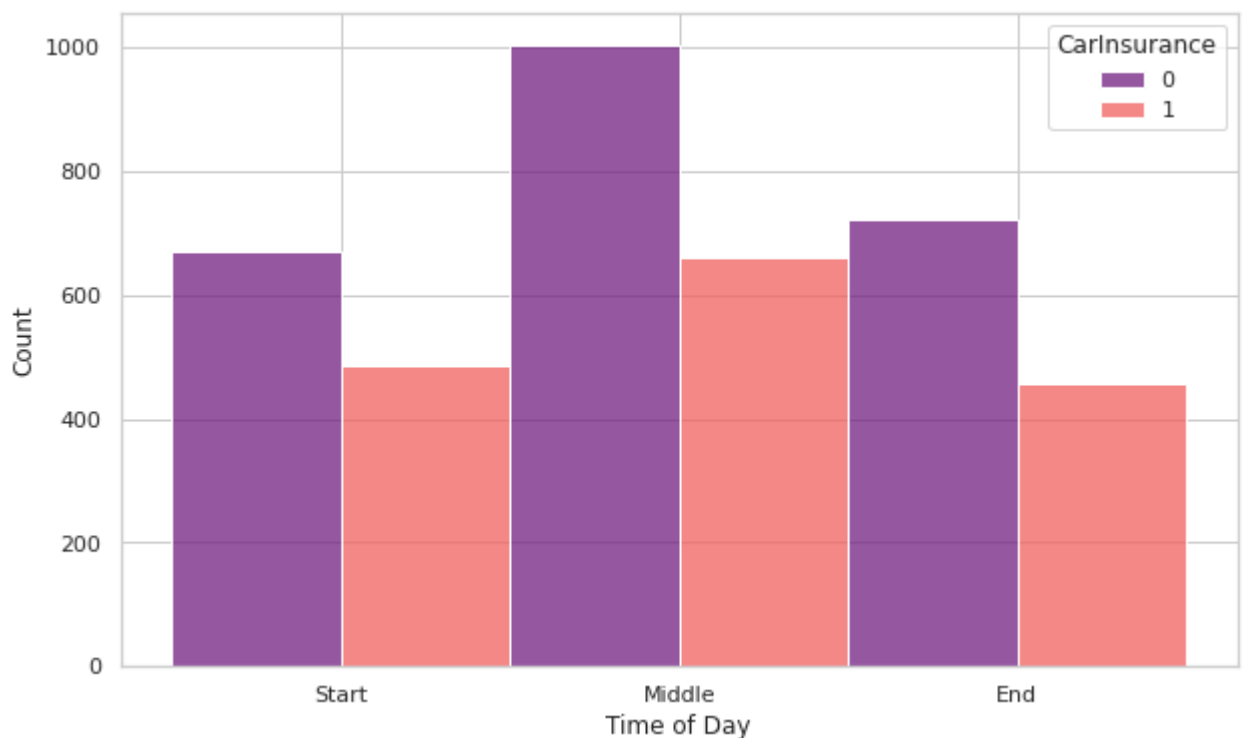
Out[71]: **Car Insurance** **0** **1** **% with Car Insurance**

DayBracket				
	Start	670	487	42.09
	Middle	1005	660	39.64
	End	721	457	38.79

In [72]: *# Plotting how the time of the month affects the likelihood of subscribing to car insurance for different age demographics*

```
plt.figure(figsize=(10,6))
sns.histplot(x='DayBracket', hue='CarInsurance', data=df_train, multiple="dodge")
plt.ylabel('Count')
plt.xlabel('Time of Day')
plt.show
```

Out[72]: <function matplotlib.pyplot.show(*args, **kw)>



In [73]: *# Get the working hours for calls*
 print(df_train["CallStart"].dt.time.min(), df_train["CallStart"].dt.time.max())
Based on the working hours (9am-6pm), I chose morning (9am-12), afternoon (12pm-6pm), and evening (6pm-9pm) day brackets.
Create the bins and give each one a label.

```
df_train['TimeBracket']=pd.cut(df_train["CallStart"].dt.hour,
                               [9,12,15,24],
                               labels=['Morning','Afternoon','Evening'],
                               right=False, include_lowest=True)
```

```
# Display the number and percentage of different contact day brackets that have
time_insured=pd.crosstab(df_train['TimeBracket'],df_train['CarInsurance'],colnam
time_insured['% with Car Insurance']=round(time_insured[1]/(time_insured[0]+time
time_insured
```

09:00:00 17:59:58


Out[73]: **Car Insurance** 0 1 % with Car Insurance

TimeBracket				
Morning	763	521		40.58
Afternoon	847	518		37.95
Evening	786	565		41.82

In [74]: *# Plotting how the time of the day affects the likelihood of subscribing to car*
for different age demographics


```
plt.figure(figsize=(10,6))
sns.histplot(x='TimeBracket', hue='CarInsurance', data=df_train, multiple="dodge")
plt.ylabel('Count')
plt.xlabel('Time of Day')
plt.show
```

Out[74]: <function matplotlib.pyplot.show(*args, **kw)>



In [75]: *# Plotting how the call duration affects the likelihood of subscribing to car in*
for different months

```
plt.figure(figsize=(10,6))
sns.barplot(x='LastContactMonth', y='CallDuration', hue='CarInsurance', data=df_
plt.ylabel('Call Duration (secs)')
plt.xlabel('Last Contact Month')
plt.show()
```

4. Prepare Data for Machine Learning

```
In [76]: #Drop columns used for exploratory analysis
df_train= df_train.drop(['Age'], axis=1)
df_train=df_train.drop(['Balance'], axis=1)
df_train=df_train.drop(['CallStart'], axis=1)
df_train=df_train.drop(['CallEnd'], axis=1)
df_train=df_train.drop(['LastContactDay'], axis=1)
df_train.head()
```

```
Out[76]:
```

	Job	Marital	Education	Default	HHInsurance	CarLoan	Communication	LastContactMont
0	management	single	tertiary	0	1	0	telephone	ja
1	blue-collar	married	primary	0	1	0	cellular	ma
2	management	single	tertiary	0	1	0	cellular	ju
3	student	single	primary	0	1	0	cellular	ma
4	management	married	tertiary	0	0	0	cellular	ju

```
In [77]: df_train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Job                    4000 non-null   object
1   Marital                4000 non-null   object
2   Education              4000 non-null   object
3   Default                4000 non-null   int64
4   HHInsurance            4000 non-null   int64
5   CarLoan                4000 non-null   int64
```

```

6   Communication      4000 non-null   object
7   LastContactMonth  4000 non-null   object
8   NoOfContacts      4000 non-null   int64
9   DaysPassed        4000 non-null   int64
10  PrevAttempts      4000 non-null   int64
11  CarInsurance      4000 non-null   int64
12  CallDuration      4000 non-null   int64
13  BalanceBracket    4000 non-null   category
14  AgeBracket        4000 non-null   category
15  DayBracket        4000 non-null   category
16  TimeBracket       4000 non-null   category
dtypes: category(4), int64(8), object(5)
memory usage: 422.8+ KB

```

Encoding categorical values

Machine learning classifiers need numerical values, so category data need to be encoded to numbers before we can fit and evaluate a model.

```

In [78]: # Label for category values
le = LabelEncoder()

df_train_cats = df_train.select_dtypes(include=['category']).columns

for i in df_train_cats:
    df_train[i] = le.fit_transform(df_train[i])

# Apply one-hot encoder to objects
df_train_objs = df_train.select_dtypes(include=['object']).columns

OH_encoder = OneHotEncoder(handle_unknown='ignore', sparse=False)
train_oh = pd.DataFrame(OH_encoder.fit_transform(df_train[df_train_objs])).astype(int)

# Get feature columns
train_oh.columns = OH_encoder.get_feature_names(df_train_objs)

# One-hot encoding removed index; put it back
train_oh.index = df_train.index

# Add one-hot encoded columns to our main df
df_train = pd.concat([df_train, train_oh], axis=1)

df_train = df_train.drop(df_train_objs, axis = 1)

```

```

In [79]: df_train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 43 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   Default                4000 non-null   int64
1   HHInsurance            4000 non-null   int64
2   CarLoan                4000 non-null   int64
3   NoOfContacts           4000 non-null   int64
4   DaysPassed             4000 non-null   int64
5   PrevAttempts           4000 non-null   int64
6   CarInsurance            4000 non-null   int64
7   CallDuration            4000 non-null   int64

```

8	BalanceBracket	4000	non-null	int64
9	AgeBracket	4000	non-null	int64
10	DayBracket	4000	non-null	int64
11	TimeBracket	4000	non-null	int64
12	Job_admin.	4000	non-null	int64
13	Job_blue-collar	4000	non-null	int64
14	Job_entrepreneur	4000	non-null	int64
15	Job_housemaid	4000	non-null	int64
16	Job_management	4000	non-null	int64
17	Job_retired	4000	non-null	int64
18	Job_self-employed	4000	non-null	int64
19	Job_services	4000	non-null	int64
20	Job_student	4000	non-null	int64
21	Job_technician	4000	non-null	int64
22	Job_unemployed	4000	non-null	int64
23	Marital_divorced	4000	non-null	int64
24	Marital_married	4000	non-null	int64
25	Marital_single	4000	non-null	int64
26	Education_primary	4000	non-null	int64
27	Education_secondary	4000	non-null	int64
28	Education_tertiary	4000	non-null	int64
29	Communication_cellular	4000	non-null	int64
30	Communication_telephone	4000	non-null	int64
31	LastContactMonth_apr	4000	non-null	int64
32	LastContactMonth_aug	4000	non-null	int64
33	LastContactMonth_dec	4000	non-null	int64
34	LastContactMonth_feb	4000	non-null	int64
35	LastContactMonth_jan	4000	non-null	int64
36	LastContactMonth_jul	4000	non-null	int64
37	LastContactMonth_jun	4000	non-null	int64
38	LastContactMonth_mar	4000	non-null	int64
39	LastContactMonth_may	4000	non-null	int64
40	LastContactMonth_nov	4000	non-null	int64
41	LastContactMonth_oct	4000	non-null	int64
42	LastContactMonth_sep	4000	non-null	int64

dtypes: int64(43)

memory usage: 1.3 MB

Update the test set to match the training set

```
In [80]: df_test.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 19 columns):

#	Column	Non-Null Count	Dtype
0	Id	1000 non-null	int64
1	Age	1000 non-null	int64
2	Job	995 non-null	object
3	Marital	1000 non-null	object
4	Education	953 non-null	object
5	Default	1000 non-null	int64
6	Balance	1000 non-null	int64
7	HHInsurance	1000 non-null	int64
8	CarLoan	1000 non-null	int64
9	Communication	779 non-null	object
10	LastContactDay	1000 non-null	int64
11	LastContactMonth	1000 non-null	object
12	NoOfContacts	1000 non-null	int64
13	DaysPassed	1000 non-null	int64
14	PrevAttempts	1000 non-null	int64
15	Outcome	243 non-null	object
16	CallStart	1000 non-null	object

```

17 CallEnd          1000 non-null    object
18 CarInsurance      0 non-null      float64
dtypes: float64(1), int64(10), object(8)
memory usage: 148.6+ KB

```

```

In [81]: # A lot of redundant code doing it this way.
# Tried concatenating the training and test sets at the beginning, but kept hav
# For time reasons I just did the test set separately here:

df_test= df_test.drop(['Id'], axis=1)
df_test= df_test.drop(['Outcome'], axis=1)

df_test[['CallStart', 'CallEnd']] = df_test[['CallStart', 'CallEnd']].astype('datetime64[ns]')

# Fill NA values with the most common attribute for each column
for column in ['Job', 'Education', 'Communication', 'Balance']:
    df_test[column] = df_test[column].fillna(df_test[column].mode()[0])

# Total Call Duration
df_test['CallDuration'] = df_test['CallEnd'] - df_test['CallStart']

# Extracting the time & converting it to seconds
df_test['CallDuration'] = df_test['CallDuration'].dt.components['minutes'] * 60 + df_test['CallDuration'].dt.components['seconds']

# Redefining the Balance column - removing the outliers
low_bal, upp_bal = remove_outliers(df_test['Balance'])
df_test['Balance'] = np.where(df_test['Balance'] > upp_bal, upp_bal, df_test['Balance'])
df_test['Balance'] = np.where(df_test['Balance'] < low_bal, low_bal, df_test['Balance'])

# Bins for Balance Bracket
test_bal_bins = [-3000, 0, 1295, 2590, 3885]
test_bal_labels = ['Negative Balance', 'Low', 'Mid', 'High']
df_test['BalanceBracket'] = pd.cut(df_test['Balance'], bins=test_bal_bins, labels=test_bal_labels)

# Bins for Age Bracket
test_age_bins = [18, 21, 31, 41, 51, 61, 71, 81, 91, 101]
test_age_labels = ['18-21', '21-30', '31-40', '41-50', '51-60', '61-70', '71-80', '81-90', '91-101']
df_test['AgeBracket'] = pd.cut(df_test['Age'], bins=test_age_bins, labels=test_age_labels)

# Bins for Day Bracket
test_day_bins = [0, 10, 21, 32]
test_day_labels = ['Start', 'Middle', 'End']
df_test['DayBracket'] = pd.cut(df_test['LastContactDay'], bins=test_day_bins, labels=test_day_labels)

# Bins for Time Bracket
df_test['TimeBracket'] = pd.cut(df_test['CallStart'].dt.hour,
                               [9, 12, 15, 24],
                               labels=['Morning', 'Afternoon', 'Evening'],
                               right=False, include_lowest=True)

# Drop columns used for exploratory analysis
df_test = df_test.drop(['Age'], axis=1)
df_test = df_test.drop(['Balance'], axis=1)
df_test = df_test.drop(['CallStart'], axis=1)
df_test = df_test.drop(['CallEnd'], axis=1)
df_test = df_test.drop(['LastContactDay'], axis=1)

df_test.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999

```

Data columns (total 17 columns):

#	Column	Non-Null Count	Dtype
0	Job	1000 non-null	object
1	Marital	1000 non-null	object
2	Education	1000 non-null	object
3	Default	1000 non-null	int64
4	HHInsurance	1000 non-null	int64
5	CarLoan	1000 non-null	int64
6	Communication	1000 non-null	object
7	LastContactMonth	1000 non-null	object
8	NoOfContacts	1000 non-null	int64
9	DaysPassed	1000 non-null	int64
10	PrevAttempts	1000 non-null	int64
11	CarInsurance	0 non-null	float64
12	CallDuration	1000 non-null	int64
13	BalanceBracket	1000 non-null	category
14	AgeBracket	1000 non-null	category
15	DayBracket	1000 non-null	category
16	TimeBracket	1000 non-null	category

dtypes: category(4), float64(1), int64(7), object(5)
memory usage: 106.4+ KB

```
In [82]: # Label the category values for Test Set
test_le = LabelEncoder()

df_test_cats = df_train.select_dtypes(include=['category']).columns

for i in df_train_cats:
    df_test[i] = test_le.fit_transform(df_test[i])

# Apply one-hot encoder to objects for Test Set
df_test_objs = df_test.select_dtypes(include=['object']).columns

OH_encoder = OneHotEncoder(handle_unknown='ignore', sparse=False)
test_oh = pd.DataFrame(OH_encoder.fit_transform(df_test[df_test_objs])).astype('float64')

# Get feature column
test_oh.columns = OH_encoder.get_feature_names(df_test_objs)

# One-hot encoding removed index; put it back
test_oh.index = df_test.index

# Add one-hot encoded columns to our main df
df_test = pd.concat([df_test, test_oh], axis=1)

df_test = df_test.drop(df_test_objs, axis = 1)

df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 43 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Default                               1000 non-null   int64
1   HHInsurance                           1000 non-null   int64
2   CarLoan                               1000 non-null   int64
3   NoOfContacts                          1000 non-null   int64
4   DaysPassed                            1000 non-null   int64
5   PrevAttempts                          1000 non-null   int64
6   CarInsurance                           0 non-null      float64
```

```

7   CallDuration      1000 non-null   int64
8   BalanceBracket    1000 non-null   int64
9   AgeBracket        1000 non-null   int64
10  DayBracket        1000 non-null   int64
11  TimeBracket       1000 non-null   int64
12  Job_admin.        1000 non-null   int64
13  Job_blue-collar   1000 non-null   int64
14  Job_entrepreneur  1000 non-null   int64
15  Job_housemaid     1000 non-null   int64
16  Job_management    1000 non-null   int64
17  Job_retired       1000 non-null   int64
18  Job_self-employed 1000 non-null   int64
19  Job_services      1000 non-null   int64
20  Job_student       1000 non-null   int64
21  Job_technician    1000 non-null   int64
22  Job_unemployed    1000 non-null   int64
23  Marital_divorced  1000 non-null   int64
24  Marital_married   1000 non-null   int64
25  Marital_single    1000 non-null   int64
26  Education_primary 1000 non-null   int64
27  Education_secondary 1000 non-null   int64
28  Education_tertiary 1000 non-null   int64
29  Communication_cellular 1000 non-null   int64
30  Communication_telephone 1000 non-null   int64
31  LastContactMonth_apr 1000 non-null   int64
32  LastContactMonth_aug 1000 non-null   int64
33  LastContactMonth_dec 1000 non-null   int64
34  LastContactMonth_feb 1000 non-null   int64
35  LastContactMonth_jan 1000 non-null   int64
36  LastContactMonth_jul 1000 non-null   int64
37  LastContactMonth_jun 1000 non-null   int64
38  LastContactMonth_mar 1000 non-null   int64
39  LastContactMonth_may 1000 non-null   int64
40  LastContactMonth_nov 1000 non-null   int64
41  LastContactMonth_oct 1000 non-null   int64
42  LastContactMonth_sep 1000 non-null   int64
dtypes: float64(1), int64(42)
memory usage: 336.1 KB

```

Split training set

X is used for the features y for CarInsurance (this is the target value we want to predict)

The Train Test split is of a 80:20 ratio respectively.

```

In [83]: X=df_train.drop(['CarInsurance'],axis=1).values
         # Including only the Target for y
         y=df_train['CarInsurance'].values

         #Splitting the Training and Testing data having 20% of Test data
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,random_

```

Define a function that will plot a confusion matrix for each classifier

5. Machine Learning Classifiers

```

In [84]: # This confusion matrix was sourced from the sklearn documentation

```

```
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    # Show all ticks
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    # Loop over data dimensions
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

class_names = ['Success', 'Failure']
```

K-Nearest Neighbour

```
In [85]: knn = KNeighborsClassifier()

# Use GridSearchCV to search over parameter values for the estimator
# This takes some time to run so can always be scaled back
parameters = {'n_neighbors':[5,6,7,8],
              'p':[1,2],
              'weights':['uniform','distance']}

knn_clf = GridSearchCV(knn, parameters)
knn_clf.fit(X_train,y_train)

# Get best parameters
print("Best parameters: ", knn_clf.best_params_, '\n')

# Accuracy
print ("kNN Accuracy: %2.2f" % accuracy_score(y_test, knn_clf.predict(X_test)))

# Evaluate score (10-fold cross validation)
score_knn = cross_val_score(knn_clf, X_train, y_train, cv=10).mean()
print("Cross Validation Score: %2.2f" % score_knn)

# Prediction
y_pred= knn_clf.predict(X_test)
print(classification_report(y_test, y_pred))

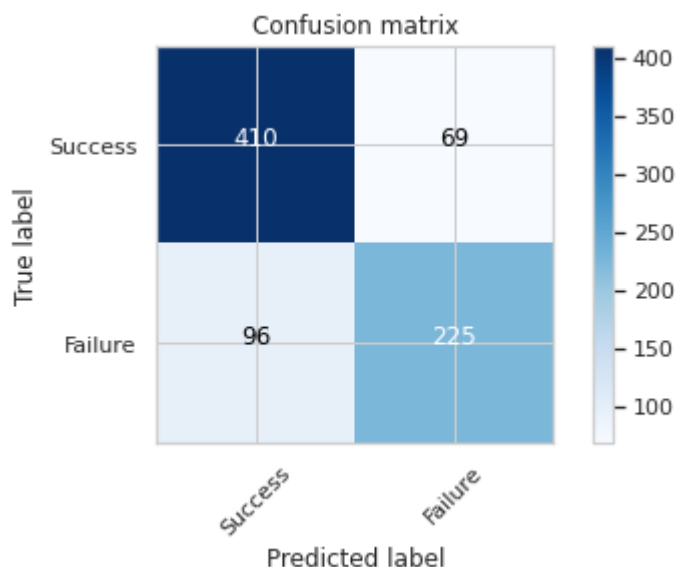
# Confusion matrix for K-Nearest Neighbour
cm = confusion_matrix(y_test,y_pred)
plot_confusion_matrix(cm, classes=class_names, title='Confusion matrix')
```

Best parameters: {'n_neighbors': 7, 'p': 1, 'weights': 'distance'}

kNN Accuracy: 0.79

Cross Validation Score: 0.77

	precision	recall	f1-score	support
0	0.81	0.86	0.83	479
1	0.77	0.70	0.73	321
accuracy			0.79	800
macro avg	0.79	0.78	0.78	800
weighted avg	0.79	0.79	0.79	800



Logistic Regression Classifier

```
In [86]: LR = LogisticRegression(max_iter=1500)

# Use GridSearchCV
parameters = {'C':[0.7,0.8,0.9,1]}

lr_clf = GridSearchCV(LR, parameters)
lr_clf.fit(X_train,y_train)

# Get best parameters
print("Best parameters: ", lr_clf.best_params_, '\n')

# Accuracy
print ("Logistic Accuracy: %2.2f" % accuracy_score(y_test, lr_clf.predict(X_test)))

# Evaluate score (10-fold cross validation)
score_LR = cross_val_score(lr_clf, X_train, y_train, cv=10).mean()
print("Cross Validation Score: %2.2f" % score_LR)

# Prediction
y_pred = lr_clf.predict(X_test)
print(classification_report(y_test, y_pred))


# Confusion matrix for Logistical Regression
cm = confusion_matrix(y_test,y_pred)
plot_confusion_matrix(cm, classes=class_names, title='Confusion matrix')
```

Best parameters: {'C': 0.8}

Logistic Accuracy: 0.81

Cross Validation Score: 0.80

	precision	recall	f1-score	support
0	0.82	0.87	0.84	479
1	0.79	0.71	0.74	321
accuracy			0.81	800
macro avg	0.80	0.79	0.79	800
weighted avg	0.80	0.81	0.80	800



Decision Tree Classifier

```
In [87]: DT = tree.DecisionTreeClassifier(random_state=21, criterion='gini')

# Use GridSearchCV
parameters = {'class_weight':[None,'balanced'],
              'min_weight_fraction_leaf':[0.0,0.01],
              'max_depth':[None,6,8,9]}

dt_clf = GridSearchCV(DT, parameters)
dt_clf.fit(X_train,y_train)

# Get best parameters
print("Best parameters: ", dt_clf.best_params_,'\n')

# Accuracy
print ("Decision Tree Accuracy: %2.2f" % accuracy_score(y_test, dt_clf.predict(y_test)))

# Evaluate score (10-fold cross validation)
score_DT = cross_val_score(dt_clf, X_train, y_train, cv=10).mean()
print("Cross Validation Score: %2.2f" % score_DT, "\n")

# Prediction
y_pred = dt_clf.predict(X_test)
print(classification_report(y_test, y_pred))

# Confusion Matrix for Decision Tree
cm = confusion_matrix(y_test,y_pred)
plot_confusion_matrix(cm, classes=class_names, title='Confusion matrix')
```

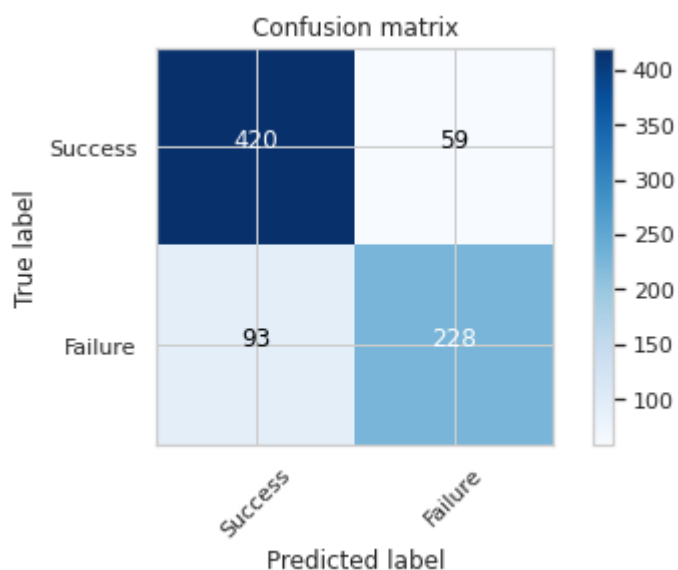
Best parameters: {'class_weight': None, 'max_depth': 6, 'min_weight_fraction_le

af': 0.01}

Decision Tree Accuracy: 0.81

Cross Validation Score: 0.79

	precision	recall	f1-score	support
0	0.82	0.88	0.85	479
1	0.79	0.71	0.75	321
accuracy			0.81	800
macro avg	0.81	0.79	0.80	800
weighted avg	0.81	0.81	0.81	800



Random Forest Classifier

```
In [88]: rfc = RandomForestClassifier(n_estimators=1000, criterion='gini', class_weight='balanced')

# Use GridSearchCV
parameters = {'min_samples_split': [8, 10, 12]}

rfc_clf = GridSearchCV(rfc, parameters)
rfc_clf.fit(X_train, y_train)

# Get best parameters
print("Best parameters: ", rfc_clf.best_params_, '\n')

# Accuracy
print("Random Forest Accuracy: %2.2f" % accuracy_score(y_test, rfc_clf.predict(X_test)))

# Evaluate score (10-fold cross validation)
score_rfc = cross_val_score(rfc_clf, X_train, y_train, cv=10).mean()
print("Cross Validation Score: %2.2f" % score_rfc)

# Prediction
y_pred = rfc_clf.predict(X_test)
print(classification_report(y_test, y_pred))

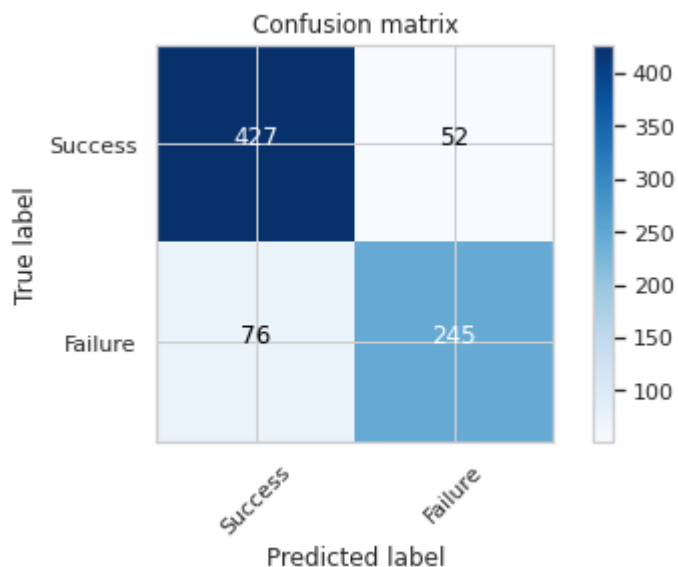
# Confusion Matrix for Random Forest
cm = confusion_matrix(y_test, y_pred)
plot_confusion_matrix(cm, classes=class_names, title='Confusion matrix')
```

Best parameters: {'min_samples_split': 12}

Random Forest Accuracy: 0.84

Cross Validation Score: 0.82

	precision	recall	f1-score	support
0	0.85	0.89	0.87	479
1	0.82	0.76	0.79	321
accuracy			0.84	800
macro avg	0.84	0.83	0.83	800
weighted avg	0.84	0.84	0.84	800



AdaBoost Classifier

```
In [43]: ada = AdaBoostClassifier(random_state=21, n_estimators=500)

# Use GridSearchCV
parameters = {'learning_rate': [0.05,0.1,0.15]}

ada_clf = GridSearchCV(ada, parameters)
ada_clf.fit(X_train,y_train)

# Get best parameters
print("Best parameters: ", ada_clf.best_params_, '\n')

# Accuracy
print ("AdaBoost Accuracy: %2.2f" % accuracy_score(y_test,ada_clf.predict(X_test)))

# Evaluate score (10-fold cross validation)
score_ada = cross_val_score(ada_clf, X_train, y_train, cv=10).mean()
print("Cross Validation Score: %2.2f" % score_ada)

# Make prediction
y_pred = ada_clf.predict(X_test)
print(classification_report(y_test,y_pred ))


# Confusion Marix for AdaBoost
cm = confusion_matrix(y_test,y_pred)
plot_confusion_matrix(cm, classes=class_names, title='Confusion matrix')
```

Best parameters: {'learning_rate': 0.15}

AdaBoost Accuracy: 0.83

Cross Validation Score: 0.82

	precision	recall	f1-score	support
0	0.84	0.88	0.86	479
1	0.81	0.75	0.78	321
accuracy			0.83	800
macro avg	0.82	0.81	0.82	800
weighted avg	0.83	0.83	0.83	800



XGBoost Classifier

```
In [89]: xgb = GradientBoostingClassifier(random_state=21, n_estimators=1000, max_depth=8)

# Use GridSearchCV
parameters = {'subsample':[0.8,0.9,1]}

xgb_clf = GridSearchCV(xgb, parameters)
xgb_clf.fit(X_train,y_train)

# Get best parameters
print("Best parameters: ", xgb_clf.best_params_, '\n')

# Accuracy
print ("GradientBoost Accuracy: %2.2f" % accuracy_score(y_test,xgb_clf.predict(y_test)))

# Evaluate score (10-fold cross validation)
score_xgb = cross_val_score(xgb_clf, X_train, y_train, cv=10).mean()
print("Cross Validation Score: %2.2f" % score_xgb)

# Prediction
y_pred = xgb_clf.predict(X_test)
print(classification_report(y_test,y_pred))


# Confusion Matrix for XGBoost
cm_xg = confusion_matrix(y_test,y_pred)
plot_confusion_matrix(cm_xg, classes=class_names, title='Confusion matrix')
```

Best parameters: {'subsample': 0.8}

GradientBoost Accuracy: 0.82

Cross Validation Score: 0.82

	precision	recall	f1-score	support
0	0.84	0.86	0.85	479
1	0.79	0.75	0.77	321
accuracy			0.82	800
macro avg	0.81	0.81	0.81	800
weighted avg	0.82	0.82	0.82	800



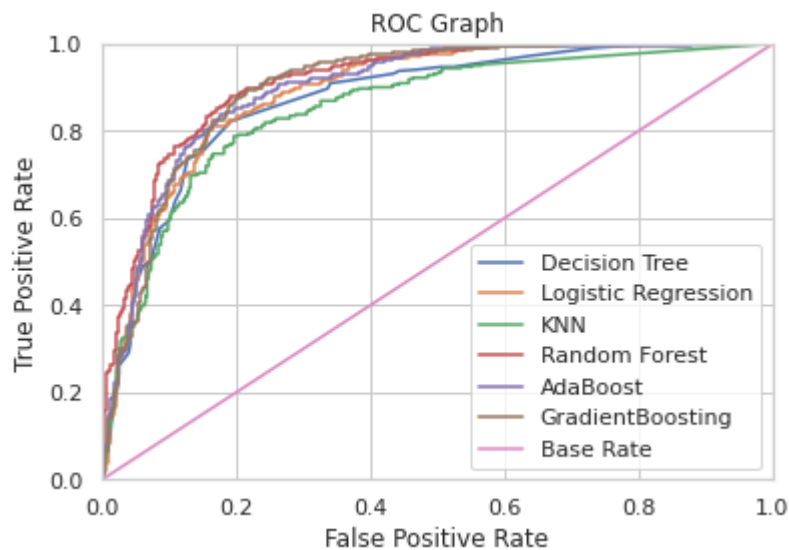
ROC Curves for each classifier

```
In [91]: # False Positive Rate, True Positive Rate and Threshold for classifiers
knn_fpr, knn_tpr, thresholds = roc_curve(y_test, knn_clf.predict_proba(X_test)[:
LR_fpr, LR_tpr, thresholds = roc_curve(y_test, lr_clf.predict_proba(X_test)[:
DT_fpr, DT_tpr, thresholds = roc_curve(y_test, dt_clf.predict_proba(X_test)[:
rfc_fpr, rfc_tpr, thresholds = roc_curve(y_test, rfc_clf.predict_proba(X_test)[:
ada_fpr, ada_tpr, thresholds = roc_curve(y_test, ada_clf.predict_proba(X_test)[:
xgb_fpr, xgb_tpr, thresholds = roc_curve(y_test, xgb_clf.predict_proba(X_test)[:

# Plotting ROC Curves for all classifiers
plt.plot(DT_fpr, DT_tpr, label='Decision Tree')
plt.plot(LR_fpr, LR_tpr, label='Logistic Regression')
plt.plot(knn_fpr, knn_tpr, label='KNN')
plt.plot(rfc_fpr, rfc_tpr, label='Random Forest')
plt.plot(ada_fpr, ada_tpr, label='AdaBoost')
plt.plot(xgb_fpr, xgb_tpr, label='GradientBoosting')

# Plot Base Rate ROC
plt.plot([0,1],[0,1],label='Base Rate')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Graph')
plt.legend(loc="lower right")
plt.show()
```



6. Testing the Trained Model

Use the best two classifiers -- Random Forest and XGBoost -- on the test data (CarInsurance_test.csv)

```
In [93]: #Random Forest Classifier
rfc.fit(X_train,y_train)
X_test = df_test.drop("CarInsurance", axis = 1)
df_test["CarInsurance"] = rfc.predict(X_test)
df_test["CarInsurance"] = df_test["CarInsurance"].apply(lambda x: "No. that will
df_test["CarInsurance"].value_counts()
```

```
Out[93]: No. that will not buy insurance [0]    602
No. that will buy insurance [1]          398
Name: CarInsurance, dtype: int64
```

```
In [94]: #XGBoost Classifier
xgb.fit(X_train,y_train)
X_test = df_test.drop("CarInsurance", axis = 1)
df_test["CarInsurance"] = xgb.predict(X_test)
df_test["CarInsurance"] = df_test["CarInsurance"].apply(lambda x: "No that will
df_test["CarInsurance"].value_counts()
```

```
Out[94]: No that will not buy insurance [0]    594
No. that will buy insurance [1]          406
Name: CarInsurance, dtype: int64
```